 [Open In Colab](#)

(<https://colab.research.google.com/github/ultralytics/yolov5/blob/master/tutorial.ipynb>)



This notebook was written by Ultralytics LLC, and is freely available for redistribution under the [GPL-3.0 license](https://choosealicense.com/licenses/gpl-3.0/) (<https://choosealicense.com/licenses/gpl-3.0/>). For more information please visit <https://github.com/ultralytics/yolov5> (<https://github.com/ultralytics/yolov5>) and <https://www.ultralytics.com> (<https://www.ultralytics.com>).

Setup

Clone repo, install dependencies, %cd into ./yolov5 folder and check GPU.

```
In [ ]: !git clone https://github.com/ultralytics/yolov5 # clone repo
!pip install -qr yolov5/requirements.txt # install dependencies (ignore errors)
%cd yolov5

import torch
from IPython.display import Image, clear_output # to display images
from utils.google_utils import download # to download models/datasets

clear_output()
print('Setup complete. Using torch %s %s' % (torch.__version__, torch.cuda.get_device_properties(0) if torch.cuda.is_available() else 'CPU'))

Setup complete. Using torch 1.5.1+cu101 _CudaDeviceProperties(name='Tesla T4', major=7, minor=5, total_memory=15079 MB, multi_processor_count=40)
```

1. Inference

Run inference with a pretrained checkpoint on contents of ./inference/images folder. Models are auto-downloaded from [Google Drive](https://drive.google.com/open?id=1Drs_Aiu7xx6S-ix95f9kNsA6ueKRpN2J) (https://drive.google.com/open?id=1Drs_Aiu7xx6S-ix95f9kNsA6ueKRpN2J).

```
In [ ]: !python detect.py --weights yolov5s.pt --img 416 --conf 0.4 --source inference/images/
Image(filename='inference/output/zidane.jpg', width=600)

Namespace(agnostic_nms=False, augment=False, classes=None, conf_thres=0.4, device='', fourcc='mp4v', half=False, img_size=416, iou_thres=0.5, output='inference/output', save_txt=False, source='./inference/images/', view_img=False, weights='yolov5s.pt')
Using CUDA device0 _CudaDeviceProperties(name='Tesla P100-PCIE-16GB', total_memory=16280MB)

image 1/2 inference/images/bus.jpg: 416x352 3 persons, 1 buss, Done. (0.009s)
image 2/2 inference/images/zidane.jpg: 288x416 2 persons, 2 ties, Done. (0.009s)
Results saved to /content/yolov5/inference/output
Done. (0.100s)
```



Inference can be run on a variety of sources: images, videos, directories, webcams, rtsp and http streams as shown in the example below.

```
In [ ]: # Example syntax (do not run cell)
!python detect.py --source file.jpg # image
        file.mp4 # video
        dir/ # directory
        0 # webcam
        'rtsp://170.93.143.139/rtp/470011e600ef003a004ee33696235daa' # rtsp
        'http://112.50.243.8/PLTV/88888888/224/3221225900/1.m3u8' # http
```

2. Test

Test a model on COCO val or test-dev dataset to determine trained accuracy. Models are auto-downloaded from [Google Drive](https://drive.google.com/open?id=1Drs_Aiu7xx6S-ix95f9kNsA6ueKRpN2J) (https://drive.google.com/open?id=1Drs_Aiu7xx6S-ix95f9kNsA6ueKRpN2J). To show results by class use the `--verbose` flag. Note that pycocotools metrics may be 1-2% better than the equivalent repo metrics, as is visible below, due to slight differences in mAP computation.

2.1 val2017

Download COCO val 2017 dataset, 1GB, 5000 images, and test model accuracy.

```
In [ ]: # Download COCO val2017
gdrive_download('1Y6Kou6kEB0ZEMCCpJSKStCor4KAReE43','coco2017val.zip') # val2017 dataset
!mv ./coco ../ # move folder alongside /yolov5
```

Downloading <https://drive.google.com/uc?export=download&id=1Y6Kou6kEB0ZEMCCpJSKStCor4KAReE43> as coco2017val.zip... unzipping... Done (11.2s)

```
In [ ]: # Run YOLOv5x on COCO val2017
!python test.py --weights yolov5x.pt --data coco.yaml --img 672
```

Namespace(augment=False, batch_size=32, conf_thres=0.001, data='./data/coco.yaml', device='', img_size=672, iou_thres=0.65, merge=False, save_json=True, single_cls=False, task='val', verbose=False, weights=['yolov5x.pt'])
Using CUDA device0_CudaDeviceProperties(name='Tesla T4', total_memory=15079MB)

Fusing layers... Model Summary: 284 layers, 8.89222e+07 parameters, 8.89222e+07 gradients
Scanning labels ../coco/labels/val2017.cache (4952 found, 0 missing, 48 empty, 0 duplicate, for 5000 images): 100% 5000/5000 [00:00<00:00, 22899.17it/s]

Class	Images	Targets	P	R	mAP@.5	mAP@.5:.95
all	5e+03	3.63e+04	0.426	0.746	0.66	0.469

Speed: 22.3/1.7/24.0 ms inference/NMS/total per 672x672 image at batch-size 32

COCO mAP with pycocotools... saving detections_val2017__results.json...

loading annotations into memory...

Done (t=0.41s)

creating index...

index created!

Loading and preparing results...

DONE (t=4.39s)

creating index...

index created!

Running per image evaluation...

Evaluate annotation type *bbox*

DONE (t=76.56s).

Accumulating evaluation results...

DONE (t=11.02s).

Average Precision (AP) @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.484

Average Precision (AP) @[IoU=0.50 | area= all | maxDets=100] = 0.668

Average Precision (AP) @[IoU=0.75 | area= all | maxDets=100] = 0.528

Average Precision (AP) @[IoU=0.50:0.95 | area= small | maxDets=100] = 0.311

Average Precision (AP) @[IoU=0.50:0.95 | area=medium | maxDets=100] = 0.534

Average Precision (AP) @[IoU=0.50:0.95 | area= large | maxDets=100] = 0.628

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets= 1] = 0.371

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets= 10] = 0.609

Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.662

Average Recall (AR) @[IoU=0.50:0.95 | area= small | maxDets=100] = 0.501

Average Recall (AR) @[IoU=0.50:0.95 | area=medium | maxDets=100] = 0.714

Average Recall (AR) @[IoU=0.50:0.95 | area= large | maxDets=100] = 0.807

2.2 test-dev2017

Download COCO test2017 dataset, 7GB, 40,000 images, to test model accuracy on test-dev set, 20,000 images. Results are saved to a *.json file which can be submitted to the evaluation server at <https://competitions.codalab.org/competitions/20794> (<https://competitions.codalab.org/competitions/20794>).

```
In [ ]: # Download COCO test-dev2017
gdrive_download('1cXZR_ckHki6nddOmcysCuuJFM--T-Q6L', 'coco2017labels.zip') # annotations
!f="test2017.zip" && curl http://images.cocodataset.org/zips/$f -o $f && unzip -q $f && rm $f # 7GB, 41k images
!mv ./test2017 ./coco/images && mv ./coco ../ # move images into /coco and move /coco alongside /yolov5

In [ ]: # Run YOLOv5s on COCO test-dev2017 with argument --task test
!python test.py --weights yolov5s.pt --data ./data/coco.yaml --task test
```

3. Train

Download <https://www.kaggle.com/ultralytics/coco128> (<https://www.kaggle.com/ultralytics/coco128>), a small 128-image tutorial dataset, start tensorboard and train YOLOv5s from a pretrained checkpoint for 3 epochs (actual training is much longer, around **300-1000 epochs**, depending on your dataset).

```
In [ ]: # Download coco128
gdrive_download('1n_oKgR81BJtk75b00eAjdv03qVCQn2f', 'coco128.zip') # coco128 dataset
!mv ./coco128 ../ # move folder alongside /yolov5
```

Downloading https://drive.google.com/uc?export=download&id=1n_oKgR81BJtk75b00eAjdv03qVCQn2f as coco128.zip... unzipping... Done (5.3s)

Train a YOLOv5s model on coco128 by specifying model config file `--cfg models/yolo5s.yaml`, and dataset config file `--data data/coco128.yaml`. Start training from pretrained `--weights yolov5s.pt`, or from randomly initialized `--weights ''`. Pretrained weights are auto-downloaded from [Google Drive](https://drive.google.com/open?id=1Drs_Aiu7xx6S-ix95f9kNsA6ueKRpN2J) (https://drive.google.com/open?id=1Drs_Aiu7xx6S-ix95f9kNsA6ueKRpN2J).

All training results are saved to `runs/exp0` for the first experiment, then `runs/exp1`, `runs/exp2` etc. for subsequent experiments.

```
In [ ]: # Start tensorboard (optional)
%load_ext tensorboard
%tensorboard --logdir runs
```

```
In [ ]: # Train YOLOv5s on coco128 for 3 epochs
```

```
python train.py --img 640 --batch 16 --epochs 3 --data coco128.yaml --cfg yolov5s.yaml --weights yolov5s.pt --nosave --cache
```

```
Namespace(batch_size=16, bucket='', cache_images=True, cfg='./models/yolov5s.yaml', data='./data/coco128.yaml', device='', epochs=3, evolve=False, hyp='', img_size=[640], multi_scale=False, name='', noautoanchor=False, nosave=True, notest=False, rect=False, resume=False, single_cls=False, weights='yolov5s.pt')
Using CUDA device0_CudaDeviceProperties(name='Tesla T4', total_memory=15079MB)
```

2020-07-11 20:37:09.422496: I tensorflow/stream_executor/platform/default/dso_loader.cc:44] Successfully opened dynamic library libcudart.so.10.1

Start Tensorboard with "tensorboard --logdir=runs", view at http://localhost:6006/

Hyperparameters {'optimizer': 'SGD', 'lr0': 0.01, 'momentum': 0.937, 'weight_decay': 0.0005, 'giou': 0.05, 'cls': 0.58, 'cls_pw': 1.0, 'obj': 1.0, 'obj_pw': 1.0, 'iou_t': 0.2, 'anchor_t': 4.0, 'fl_gamma': 0.0, 'hsv_h': 0.014, 'hsv_s': 0.68, 'hsv_v': 0.36, 'degrees': 0.0, 'translate': 0.0, 'scale': 0.5, 'shear': 0.0}

	from	n	params	module	arguments
0	-1	1	3520	models.common.Focus	[3, 32, 3]
1	-1	1	18560	models.common.Conv	[32, 64, 3, 2]
2	-1	1	19904	models.common.BottleneckCSP	[64, 64, 1]
3	-1	1	73984	models.common.Conv	[64, 128, 3, 2]
4	-1	1	161152	models.common.BottleneckCSP	[128, 128, 3]
5	-1	1	295424	models.common.Conv	[128, 256, 3, 2]
6	-1	1	641792	models.common.BottleneckCSP	[256, 256, 3]
7	-1	1	1180672	models.common.Conv	[256, 512, 3, 2]
8	-1	1	656896	models.common.SPP	[512, 512, [5, 9, 13]]
9	-1	1	1248768	models.common.BottleneckCSP	[512, 512, 1, False]
10	-1	1	131584	models.common.Conv	[512, 256, 1, 1]
11	-1	1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
12	[-1, 6]	1	0	models.common.Concat	[1]
13	-1	1	378624	models.common.BottleneckCSP	[512, 256, 1, False]
14	-1	1	33024	models.common.Conv	[256, 128, 1, 1]
15	-1	1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
16	[-1, 4]	1	0	models.common.Concat	[1]
17	-1	1	95104	models.common.BottleneckCSP	[256, 128, 1, False]
18	-1	1	32895	torch.nn.modules.conv.Conv2d	[128, 255, 1, 1]
19	-2	1	147712	models.common.Conv	[128, 128, 3, 2]
20	[-1, 14]	1	0	models.common.Concat	[1]
21	-1	1	313088	models.common.BottleneckCSP	[256, 256, 1, False]
22	-1	1	65535	torch.nn.modules.conv.Conv2d	[256, 255, 1, 1]
23	-2	1	590336	models.common.Conv	[256, 256, 3, 2]
24	[-1, 10]	1	0	models.common.Concat	[1]
25	-1	1	1248768	models.common.BottleneckCSP	[512, 512, 1, False]
26	-1	1	130815	torch.nn.modules.conv.Conv2d	[512, 255, 1, 1]
27	[-1, 22, 18]	1	0	models.yolo.Detect	[80, [[116, 90, 156, 198, 373, 326], [30, 61, 62, 45, 59, 119], [10, 13, 16, 30, 33, 23]]]

Model Summary: 191 layers, 7.46816e+06 parameters, 7.46816e+06 gradients

Optimizer groups: 62 .bias, 70 conv.weight, 59 other

Scanning labels ./coco128/labels/train2017.cache (126 found, 0 missing, 2 empty, 0 duplicate, for 128 images): 100% 128/128 [00:00<00:00, 20484.22it/s]

Caching images (0.1GB): 100% 128/128 [00:00<00:00, 156.07it/s]

Scanning labels ./coco128/labels/train2017.cache (126 found, 0 missing, 2 empty, 0 duplicate, for 128 images): 100% 128/128 [00:00<00:00, 22082.55it/s]

Caching images (0.1GB): 100% 128/128 [00:00<00:00, 152.91it/s]

Analyzing anchors... Best Possible Recall (BPR) = 0.9935

Image sizes 640 train, 640 test

Using 2 dataloader workers

Starting training for 3 epochs...

Epoch	gpu_mem	GloU	obj	cls	total	targets	img_size
0/2	6.84G	0.04376	0.06831	0.02	0.1321	225	640: 100% 8/8 [00:09<00:00, 1.22s/it]
Class	Images	Targets	P	R	mAP@.5	mAP@.5:.95	100% 8/8 [00:09<00:00, 1.24s/it]
all	128	929	0.34	0.762	0.69	0.446	

Epoch	gpu_mem	GloU	obj	cls	total	targets	img_size
1/2	6.06G	0.04333	0.08225	0.02207	0.1476	182	640: 100% 8/8 [00:03<00:00, 2.17it/s]
Class	Images	Targets	P	R	mAP@.5	mAP@.5:.95	100% 8/8 [00:02<00:00, 3.28it/s]
all	128	929	0.342	0.755	0.687	0.447	

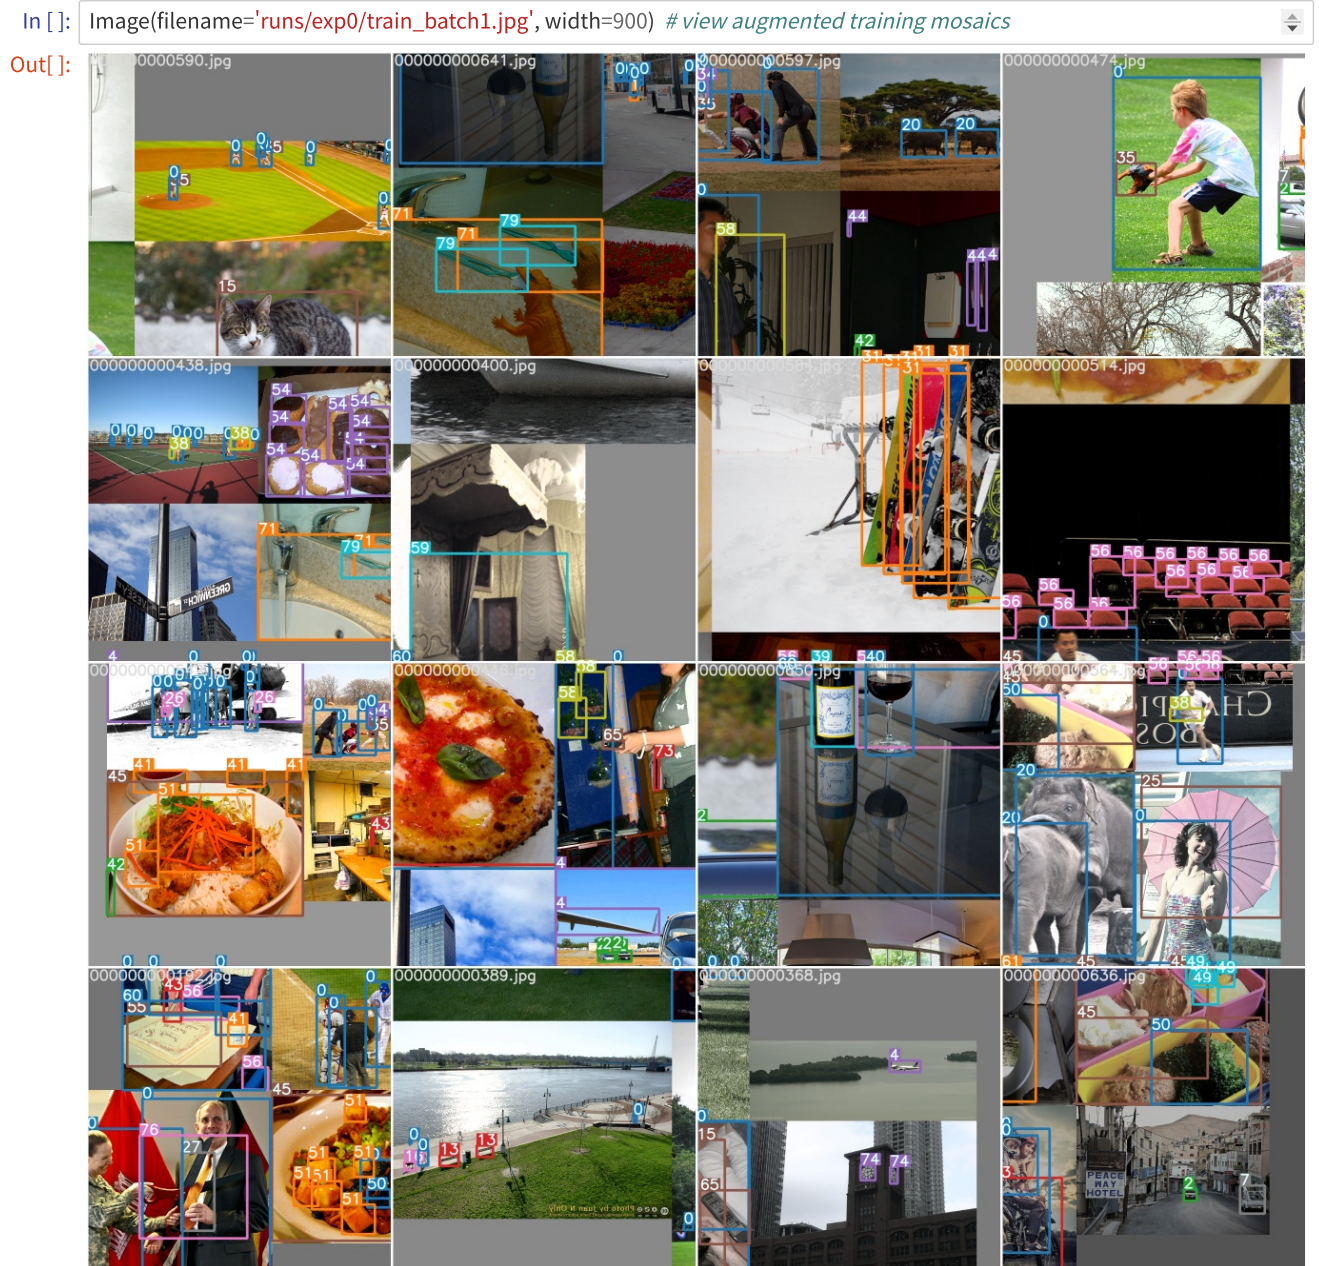
Epoch	gpu_mem	GloU	obj	cls	total	targets	img_size
2/2	6.06G	0.0444	0.07251	0.01855	0.1355	216	640: 100% 8/8 [00:03<00:00, 2.15it/s]
Class	Images	Targets	P	R	mAP@.5	mAP@.5:.95	100% 8/8 [00:02<00:00, 3.46it/s]
all	128	929	0.354	0.759	0.689	0.45	

Optimizer stripped from runs/exp0/weights/last.pt, 15.2MB

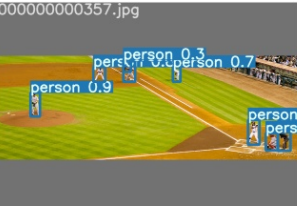

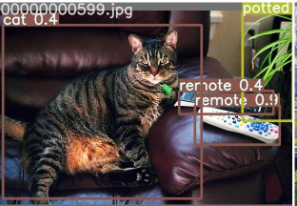
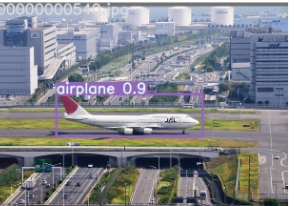

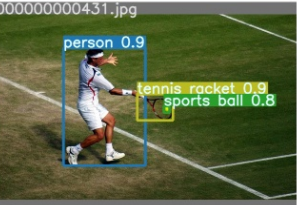

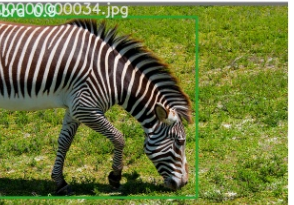







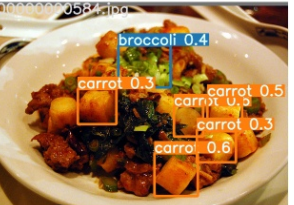
3 epochs completed in 0.009 hours.

4. Visualize

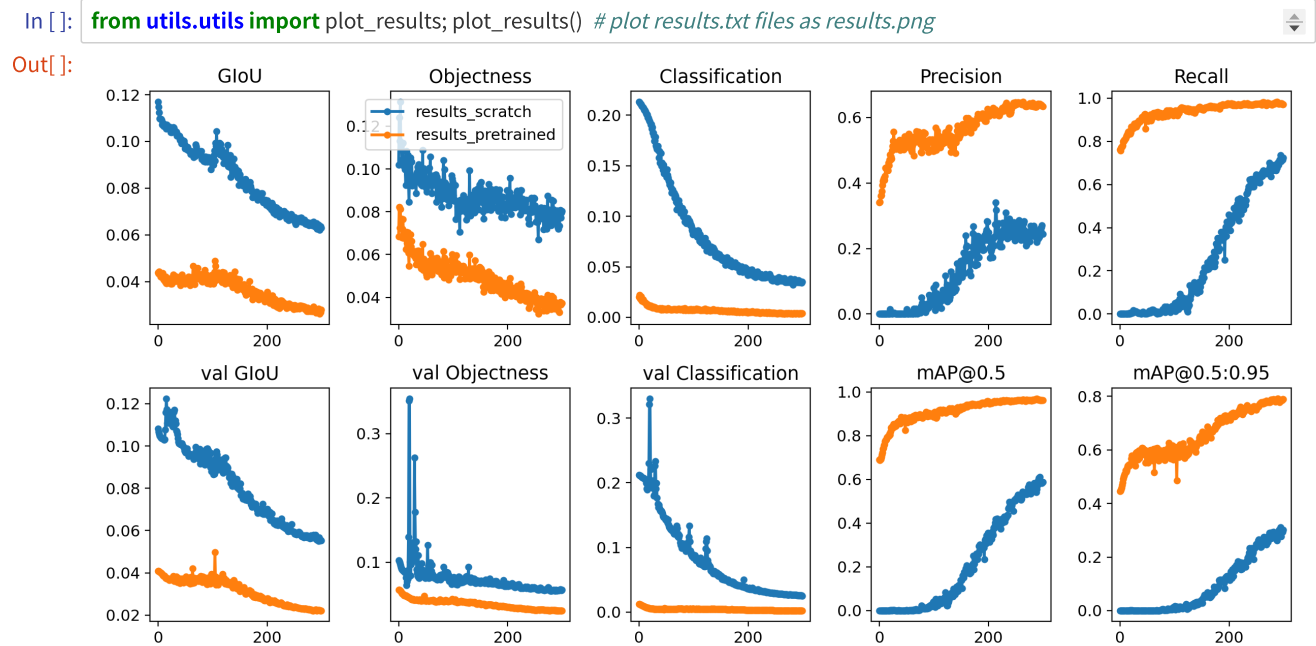
View `runs/exp0/train*.jpg` images to see training images, labels and augmentation effects. A **Mosaic Dataloader** is used for training (shown below), a new concept developed by Ultralytics and first featured in [YOLOv4](https://arxiv.org/abs/2004.10934) (<https://arxiv.org/abs/2004.10934>).



View `test_batch0_gt.jpg` to see test batch 0 *ground truth* labels.

000000000357.jpg  person 0.3 person 0.7 person 0.9 person 0.9	000000000428.jpg  person 0.9 cake 0.8	000000000599.jpg  cat 0.4 remote 0.4 remote 0.4	000000000514.jpg  airplane 0.9
000000000472.jpg  airplane 0.9	000000000431.jpg  person 0.9 tennis racket 0.9 sports ball 0.8	000000000599.jpg  person 0.3 person 0.5 person 0.5 bench 0.3	000000000034.jpg  zebra 0.9
000000000597.jpg  elephant 0.8 elephant 0.8 elephant 0.9 elephant 0.9	000000000491.jpg  teddy bear 0.4 teddy bear 0.5 teddy bear 0.3	00000000047.jpg  airplane 0.5 car 0.8	00000000081.jpg  airplane 0.7
000000000338.jpg  clock 0.9 refrigerator 0.7 person 0.7 person 0.8 sink 0.5 chair 0.8	000000000458.jpg  person 0.9 person 0.5 baseball bat 0.3 person 0.9 baseball glove 0.3 sports ball 0.9	000000000359.jpg  traffic light 0.6 traffic light 0.6 car 0.5	000000000584.jpg  broccoli 0.4 carrot 0.3 carrot 0.5 carrot 0.3 carrot 0.6

Training losses and performance metrics are saved to Tensorboard and also to a `runs/exp0/results.txt` logfile. `results.txt` is plotted as `results.png` after training completes. Partially completed `results.txt` files can be plotted with `from utils.utils import plot_results; plot_results()`. Here we show YOLOv5s trained on coco128 to 300 epochs, starting from scratch (blue), and from pretrained `yolov5s.pt` (orange).



Environments

YOLOv5 may be run in any of the following up-to-date verified environments (with all dependencies including CUDA/CUDNN, Python and PyTorch preinstalled):

- **Google Colab Notebook** with free GPU: [Open In Colab](https://colab.research.google.com/github/ultralytics/yolov5/blob/master/tutorial.ipynb) (<https://colab.research.google.com/github/ultralytics/yolov5/blob/master/tutorial.ipynb>)
- **Kaggle Notebook** with free GPU: <https://www.kaggle.com/ultralytics/yolov5> (<https://www.kaggle.com/ultralytics/yolov5>)
- **Google Cloud** Deep Learning VM. See [GCP Quickstart Guide](https://github.com/ultralytics/yolov5/wiki/GCP-Quickstart) (<https://github.com/ultralytics/yolov5/wiki/GCP-Quickstart>)
- **Docker Image** <https://hub.docker.com/r/ultralytics/yolov5> (<https://hub.docker.com/r/ultralytics/yolov5>). See [Docker Quickstart Guide](https://github.com/ultralytics/yolov5/wiki/Docker-Quickstart) (<https://github.com/ultralytics/yolov5/wiki/Docker-Quickstart>)



Appendix

Optional extras below. Unit tests validate repo functionality and should be run on any PRs submitted.

```
In [ ]: # Re-clone repo
%cd ..
!rm -rf yolov5 && git clone https://github.com/ultralytics/yolov5
%cd yolov5
```

```
In [ ]: # Test GCP ckpt
%%shell
for x in best*
do
  gsutil cp gs://*/*/$x.pt .
  python test.py --weights $x.pt --data coco.yaml --img 672
done
```

```
In [ ]: # YOLOv5 unit tests
%%shell
cd .. && rm -rf yolov5 && git clone https://github.com/ultralytics/yolov5 && cd yolov5
export PYTHONPATH="$PWD" # to run *.py. files in subdirectories
pip install -qr requirements.txt onnx
python3 -c "from utils.google_utils import *; gdrive_download('1n_oKgR81BJtk75b00eAjdv03qVCQn2f', 'coco128.zip')" &
& mv ./coco128 ../
for x in yolov5s yolov5m yolov5l yolov5x # models
do
python train.py --weights $x.pt --cfg $x.yaml --epochs 4 --img 320 --device 0 # train
for di in 0 cpu # inference devices
do
python detect.py --weights $x.pt --device $di # detect official
python detect.py --weights runs/exp0/weights/last.pt --device $di # detect custom
python test.py --weights $x.pt --device $di # test official
python test.py --weights runs/exp0/weights/last.pt --device $di # test custom
done
python models/yolo.py --cfg $x.yaml # inspect
python models/export.py --weights $x.pt --img 640 --batch 1 # export
done
```